

CUDA graph break-even on H100: launch overhead vs. iterations

